

# Bayesian Statistical Methods

## Partial solutions

### Chapter 3: Computational approaches

Jump to problem: [1](#), [3](#), [5](#), [7](#), [9](#), [11](#), [13](#), [15](#), [17](#)

(1) MAP: Fast but doesn't quantify uncertainty

Numerical integration: Accurate but fails in high dimensions

CLT: Fast but requires a large sample size

Gibbs: Accurate but requires conjugate priors

M-H: Flexible but requires tuning

(3a) The conjugate prior is  $\lambda \sim \text{Gamma}(a, b)$  and the posterior is

$$\begin{aligned} p(\lambda | Y) &\propto \left[ \prod_{i=1}^n f(Y_i | \lambda) \right] \pi(\lambda) \\ &\propto \left[ \prod_{i=1}^n \lambda^{Y_i} \exp(-N_i \lambda) \right] \cdot [\lambda^{a-1} \exp(-\lambda b)] \\ &\propto \lambda^{(\sum_{i=1}^n Y_i + a) - 1} \exp \left[ -\lambda \left( \sum_{i=1}^n N_i + b \right) \right] \end{aligned}$$

and so  $\lambda | Y \sim \text{Gamma}(\sum_{i=1}^n Y_i + a, \sum_{i=1}^n N_i + b)$ .

(3b) For  $\lambda \in (0, 20)$  the posterior is  $p(\lambda | Y) \propto \prod_{i=1}^n f(Y_i | \lambda)$ . Therefore the log posterior is a constant plus

$$l(\lambda) = \sum_{i=1}^n [-N_i \lambda + Y_i \log(\lambda)].$$

Taking the derivative and setting to zero gives

$$l'(\lambda) = -\sum_{i=1}^n N_i + \left[ \sum_{i=1}^n Y_i \right] / \lambda = 0$$

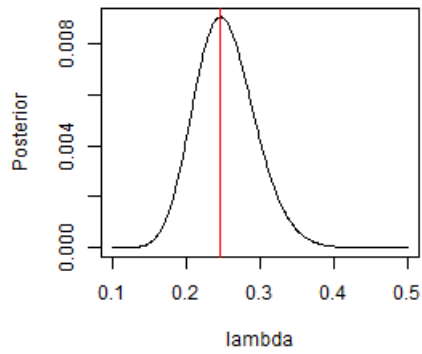
and so the MAP estimate is  $\hat{\lambda} = [\sum_{i=1}^n Y_i] / [\sum_{i=1}^n N_i]$  (unless this is greater than 20, in which case the MAP is 20).

(3c)

```
lambda <- seq(0.1,0.5,0.001)
post    <- dpois(12,50*lambda)*dpois(25,100*lambda)
MAP     <- (12+25)/(50+100)
MAP
```

```
## [1] 0.2466667
```

```
plot(lambda,post,type="l",ylab="Posterior")
abline(v=MAP,col=2)
```

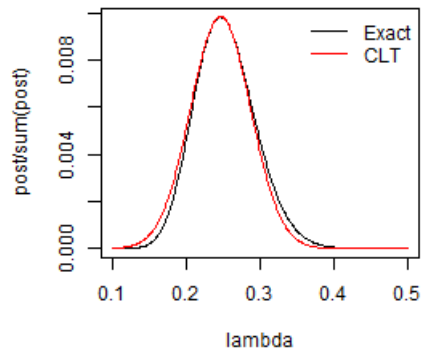


(3d) Taking the second derivative gives

$$\frac{d^2 l(\lambda)}{d\lambda^2} = -\sum_{i=1}^n \frac{Y_i}{\lambda^2}$$

and so the approximate variance is  $\hat{\lambda}^2 / [\sum_{i=1}^n Y_i] = [\sum_{i=1}^n Y_i] / [\sum_{i=1}^n N_i]^2$ .

```
lambda <- seq(0.1,0.5,0.001)
MAP <- (12+25)/(50+100)
VAR <- MAP*MAP/(12+25)
clt <- dnorm(lambda,MAP,sqrt(VAR))
plot(lambda,post/sum(post),type="l")
lines(lambda,clt/sum(clt),col=2)
legend("topright",c("Exact","CLT"),lty=1,col=1:2,bty="n")
```



(5a) Perhaps  $Y_i$  is the number of crimes in year  $i$  and a new police policy is put in place in year  $n$  so the objective is to test whether the mean number of crimes is affected by the new policy.

(5b) Let  $\tau = 1/\sigma^2$ .

$$\mu | \text{rest} \sim \text{Normal} \left( \frac{\tau \sum_{i=1}^N Y_i + \tau \sum_{i=n+1}^{n+m} (Y_i - \delta)}{\tau(n+m) + 1/100^2}, \frac{1}{\tau(n+m) + 1/100^2} \right).$$

Similarly,

$$\delta | \text{rest} \sim \text{Normal} \left( \frac{\tau \sum_{i=n+1}^{n+m} (Y_i - \mu)}{\tau m + 1/100^2}, \frac{1}{\tau m + 1/100^2} \right).$$

Finally,

$$\sigma^2 | \text{rest} \sim \text{InvGamma} \left( a + (n+m)/2, b + \sum_{i=1}^n (Y_i - \mu)^2/2 + \sum_{i=n+1}^{n+m} (Y_i - \mu - \delta)^2/2 \right).$$

(5c)

```
# Generate data
n      <- 50
m      <- 50
mu_true <- 10
delta_true <- 1
sigma_true <- 2
set.seed(919)
Y1     <- rnorm(n,mu_true,sigma_true)
Y2     <- rnorm(m,mu_true+delta_true,sigma_true)
Y      <- c(Y1,Y2)

# Prep for Gibbs sampling
S      <- 10000 # number of iterations
pri_var <- 100^2 # priors
eps    <- 0.01
keep   <- matrix(0,S,3)
mu     <- mean(Y) # initial values
delta  <- 0
sigma2 <- var(Y)

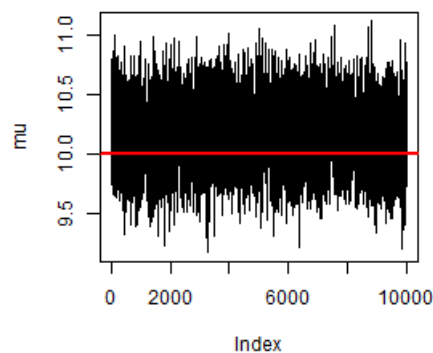
# Go!
for(iter in 1:S){
  # mu
  prec <- (n+m)/sigma2 + 1/pri_var
  mn   <- sum(Y[1:n])/sigma2 + sum(Y[1:m+n]-delta)/sigma2
  mu   <- rnorm(1,mn/prec,1/sqrt(prec))

  # delta
  prec <- m/sigma2 + 1/pri_var
  mn   <- sum(Y[1:m+n]-mu)/sigma2
  delta <- rnorm(1,mn/prec,1/sqrt(prec))

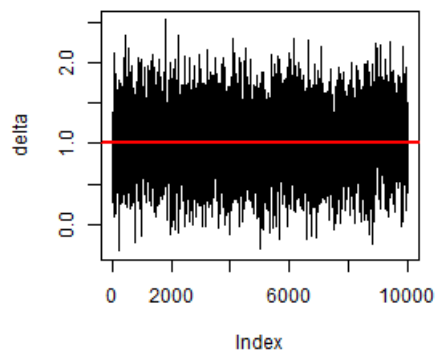
  # sigma2
  SS   <- sum((Y[1:n]-mu)^2) +
          sum((Y[1:m+n]-mu-delta)^2)
  sigma2 <- 1/rgamma(1,eps+(n+m)/2,SS/2+eps)

  keep[iter,] <- c(mu,delta,sigma2)
}

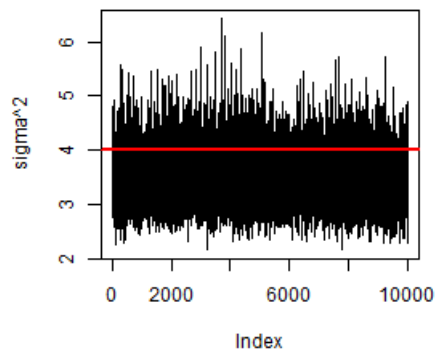
plot(keep[,1],type="l",ylab="mu")
abline(mu_true,0,lwd=2,col=2)
```



```
plot(keep[,2],type="l",ylab="delta")
abline(delta_true,0,lwd=2,col=2)
```



```
plot(keep[,3],type="l",ylab="sigma^2")
abline(sigma_true^2,0,lwd=2,col=2)
```



(5d)

```

library(rjags)
data <- list(n=n,m=m,Y1=Y[1:n],Y2=Y[n+1:m])

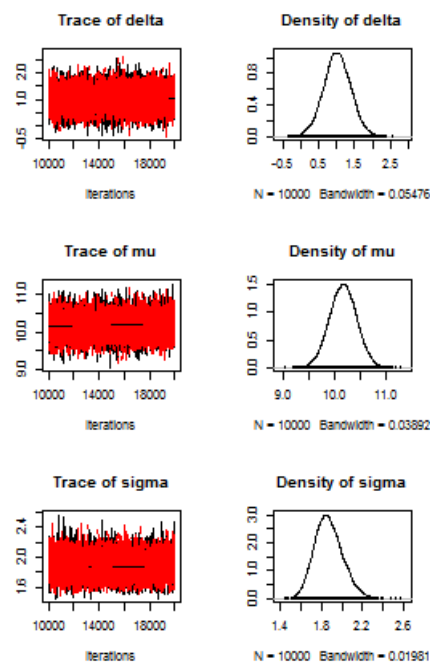
model_string <- textConnection("model{

# Likelihood
for(i in 1:n){
  Y1[i] ~ dnorm(mu,tau)
}
for(i in 1:m){
  Y2[i] ~ dnorm(mu+delta,tau)
}

# Priors
mu ~ dnorm(0, 0.0001)
delta ~ dnorm(0, 0.0001)
tau ~ dgamma(0.1, 0.1)
sigma <- 1/sqrt(tau)
}")

inits <- list(mu=mean(Y),delta=0,tau=1/var(Y))
model <- jags.model(model_string,data = data, inits=inits, n.chains=2,quiet=TRUE)
update(model, 10000, progress.bar="none")
params <- c("mu","delta","sigma")
samples <- coda.samples(model,
  variable.names=params,
  n.iter=10000, progress.bar="none")
plot(samples)

```



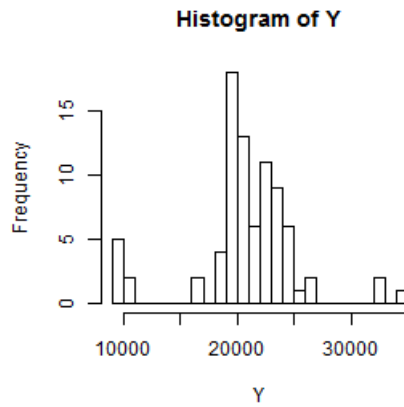
Convergence looks great and the posterior approximations in parts c and d are very similar.

(7a)

```

library(MASS)
Y<-galaxies
n<-length(Y)
hist(Y,breaks=25)

```



```
mean(Y);var(Y)
```

```
## [1] 20828.17
```

```
## [1] 20827887
```

One reasonable initial value is to start at the approximately Gaussian model with  $k = 30$  and mean and variance set to the sample mean and variance (above).

(7b)

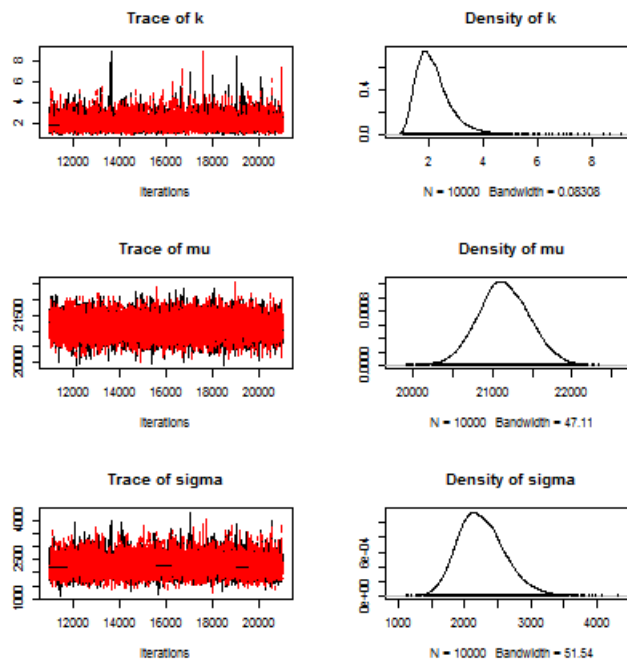
```
library(rjags)
data <- list(n=n,Y=Y)

model_string <- textConnection("model{

# Likelihood
for(i in 1:n){
  Y[i] ~ dt(mu,tau,k)
}

# Priors
mu ~ dnorm(0, 0.0000000001)
tau ~ dgamma(0.01, 0.01)
k ~ dunif(1, 30)
sigma <- 1/sqrt(tau)
}")

inits <- list(mu=mean(Y),tau=1/var(Y),k=30)
model <- jags.model(model_string,data = data, inits=inits, n.chains=2,quiet=TRUE)
update(model, 10000, progress.bar="none")
params <- c("mu","sigma","k")
samples <- coda.samples(model,
  variable.names=params,
  n.iter=10000, progress.bar="none")
plot(samples)
```



```
summary(samples)
```

```
##
## Iterations = 11001:21000
## Thinning interval = 1
## Number of chains = 2
## Sample size per chain = 10000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean      SD Naive SE Time-series SE
## k           2.158    0.682 0.004823      0.0111
## mu        21140.856 323.627 2.288385      2.9208
## sigma     2262.255 359.784 2.544059      4.9864
##
## 2. Quantiles for each variable:
##
##           2.5%      25%      50%      75%      97.5%
## k           1.247    1.701    2.034    2.462    3.792
## mu        20515.753 20925.426 21133.193 21357.048 21788.130
## sigma     1645.335  2008.107  2230.688  2480.372  3053.795
```

```
mn <- summary(samples)$statistics[,1]
mn
```

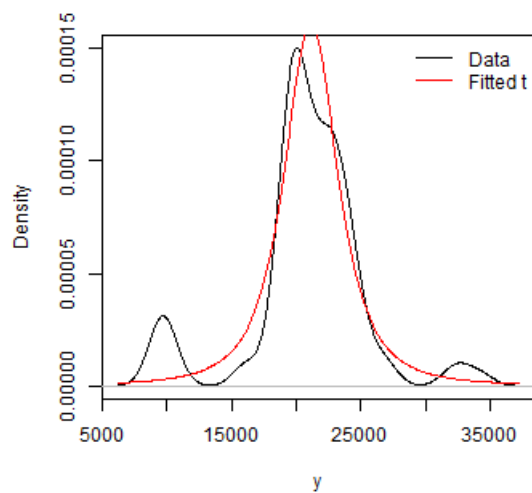
```
##           k           mu           sigma
##    2.158073 21140.856281  2262.254839
```

Convergence looks good.

(7c) The code below compares the fitted t PDF with the kernel density estimate (an alternative to the histogram) of the data.

```
kde <- density(Y)
pdf <- dt((kde$x-mn[2])/mn[3],df=mn[1])
pdf <- sum(kde$y)*pdf/sum(pdf) # This makes both plots on the same scale

plot(kde,xlab="y",ylab="Density",main="")
lines(kde$x,pdf,col=2)
legend("topright",c("Data", "Fitted t"),lty=1,col=1:2,bty="n")
```



The t density misses the humps on the left and right and isn't a great fit.

(9)

```
library(rjags)
data <- list(N1=2820,Y1=563,N2=27,Y2=10)

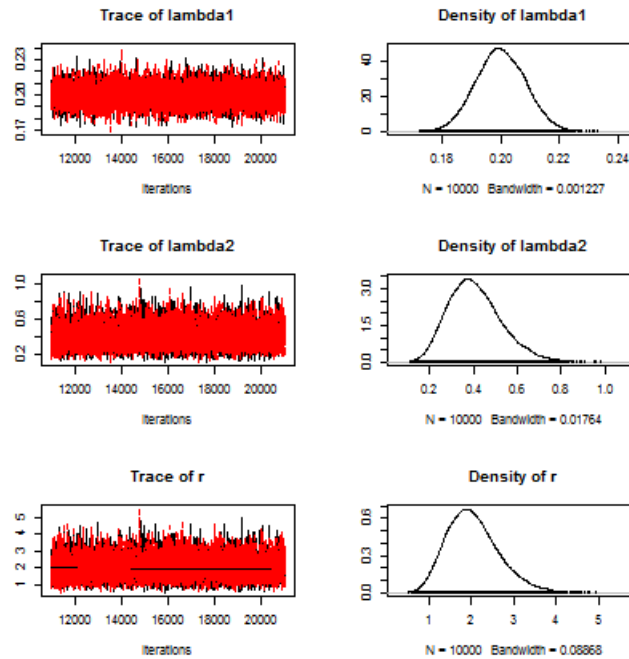
model_string <- textConnection("model{

# Likelihood
Y1 ~ dpois(N1*lambda1)
Y2 ~ dpois(N2*lambda2)

# Priors
lambda1 ~ dunif(0, 10)
lambda2 ~ dunif(0, 10)
r      <- lambda2/lambda1
}")

inits <- list(lambda1=0.1,lambda2=0.2)
model <- jags.model(model_string,data = data, inits=inits, n.chains=2,quiet=TRUE)
update(model, 10000, progress.bar="none")
params <- c("lambda1","lambda2","r")
samples <- coda.samples(model,
  variable.names=params,
  n.iter=10000, progress.bar="none")
plot(samples)
```





```
summary(samples)
```

```
##
## Iterations = 11001:21000
## Thinning interval = 1
## Number of chains = 2
## Sample size per chain = 10000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean      SD Naive SE Time-series SE
## lambda1 0.2000 0.008387 5.931e-05 7.531e-05
## lambda2 0.4079 0.122202 8.641e-04 1.148e-03
## r       2.0428 0.617442 4.366e-03 5.730e-03
##
## 2. Quantiles for each variable:
##
##           2.5%  25%   50%   75%  97.5%
## lambda1 0.1839 0.1943 0.1998 0.2057 0.2166
## lambda2 0.2044 0.3200 0.3953 0.4816 0.6819
## r       1.0171 1.6016 1.9809 2.4141 3.4344
```

```
effectiveSize(samples)
```

```
## lambda1 lambda2      r
## 12402.53 11332.00 11627.71
```

```
gelman.diag(samples)
```

```
## Potential scale reduction factors:
##
##           Point est. Upper C.I.
## lambda1         1         1
## lambda2         1         1
## r               1         1
##
## Multivariate psrf
##
## 1
```

The trace plots look great, the effective sample sizes are all large (over 1,000), and the Gelman-Rubin statistics are 1.0. Therefore, the chains have clearly converged.

(11) The code below uses Gaussian candidate distributions with standard deviation 0.2 and 0.05 (selected by trial and error to give acceptance probability around 0.4 for each parameter).

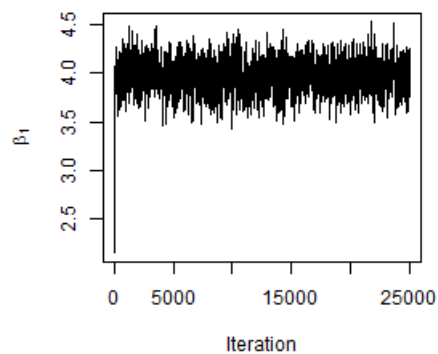
```
Y      <- c(64,13,33,18,30,20)
t      <- 1:6

S      <- 25000
beta   <- c(2,0)
cansd  <- c(0.2,0.05)
keep   <- matrix(0,S,2)

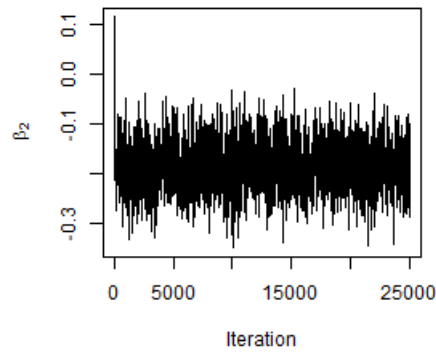
post   <- function(Y,t,beta,pri.sd=10){
  mn <- exp(beta[1] + t*beta[2])
  l   <- prod(dpois(Y,mn))
  p   <- prod(dnorm(beta,0,pri.sd))
  return(l*p)}

for(iter in 1:S){
  for(j in 1:2){
    can  <- beta
    can[j] <- rnorm(1,beta[j],cansd[j])
    R     <- post(Y,t,can)/post(Y,t,beta)
    if(runif(1)<R){
      beta <- can
    }
  }
  keep[iter,] <- beta
}

plot(keep[,1],type="l",ylab=expression(beta[1]),xlab="Iteration")
```



```
plot(keep[,2],type="l",ylab=expression(beta[2]),xlab="Iteration")
```



```
acc_rate <- colMeans(keep[-1,]!=keep[-S,])
acc_rate
```

```
## [1] 0.4033761 0.4568983
```

The trace plots look great and the acceptance rates are around 0.4, which is acceptable.

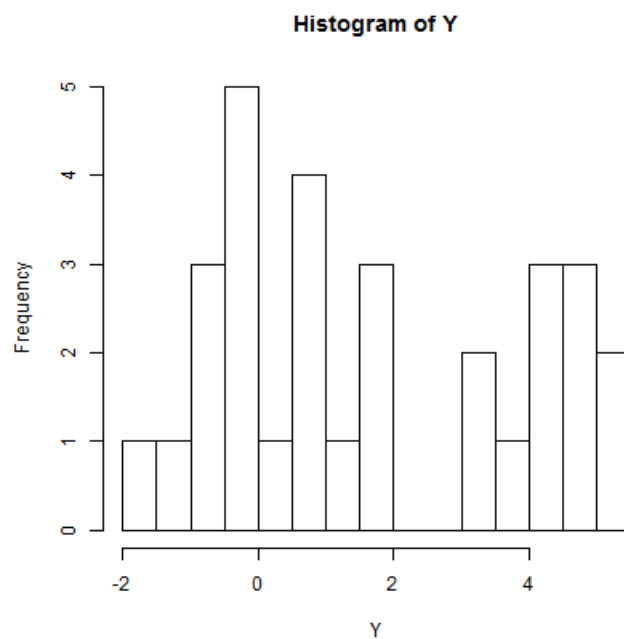
(13a) The joint distribution of  $B$  and  $Y$  is  $f(Y, B) = g(Y|B)h(B)$  where  $g(Y|B) = \phi(Y - B\theta)$  is the normal PDF and  $h$  is the Bernoulli PMF. The marginal distribution of  $Y$  is then

$$f(Y) = \sum_{b=0}^1 g(Y|B)h(B) = \phi(Y)^{\frac{1}{2}} + \phi(Y-\theta)^{\frac{1}{2}}$$

as desired.

(13b)

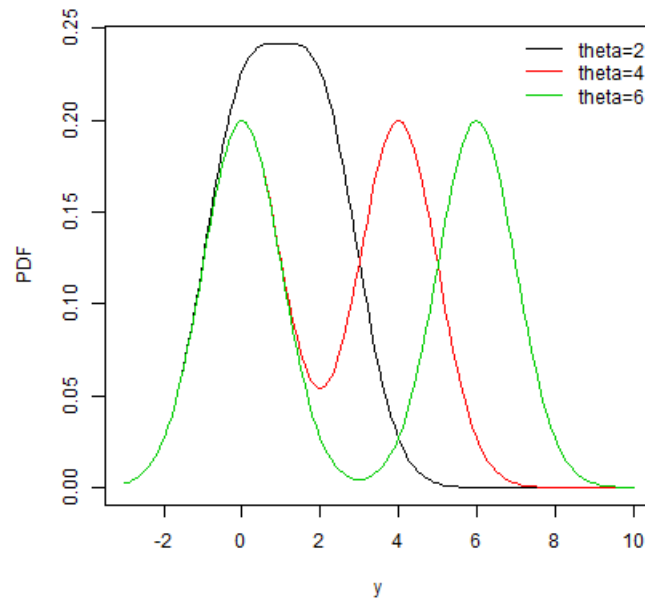
```
set.seed(27695)
theta_true <- 4
n          <- 30
B          <- rbinom(n,1,0.5)
Y          <- rnorm(n,B*theta_true,1)
hist(Y,breaks=25)
```



```

y <- seq(-3,10,0.1)
plot(y,0.5*dnorm(y,0,1) + 0.5*dnorm(y,2,1),type="l",ylab="PDF")
lines(y,0.5*dnorm(y,0,1) + 0.5*dnorm(y,4,1),col=2)
lines(y,0.5*dnorm(y,0,1) + 0.5*dnorm(y,6,1),col=3)
legend("topright",c("theta=2","theta=4","theta=6"),
      lty=1,col=1:3,bty="n")

```



(13c)

```

library(stats4)
nlp <- function(theta,Y){
  like <- 0.5*dnorm(Y,0,1)+
    0.5*dnorm(Y,theta,1)
  prior <- dnorm(theta,0,10)
  neg_log_post <- -sum(log(like))-log(prior)
  return(neg_log_post)}

map_est <- mle(nlp,start=list(theta=1),
              fixed=list(Y=Y))
sd <- sqrt(vcov(map_est))
map_est; sd

```

```

##
## Call:
## mle(minuslogl = nlp, start = list(theta = 1), fixed = list(Y = Y))
##
## Coefficients:
##      theta      Y1      Y2      Y3      Y4      Y5
## 4.17919603 -1.77844991 1.60305490 0.63529137 4.67157916 -0.43544347
##      Y6      Y7      Y8      Y9     Y10     Y11
## -0.12584690 -0.38585384 4.66977750 1.12469067 1.79842288 3.85616261
##     Y12     Y13     Y14     Y15     Y16     Y17
## 3.12904919 -1.38070620 4.07364670 3.35284203 0.52304572 -0.75516014
##     Y18     Y19     Y20     Y21     Y22     Y23
## -0.52569447 0.76633658 5.04733582 0.11272642 4.78594654 1.69632947
##     Y24     Y25     Y26     Y27     Y28     Y29
## 4.05891960 0.58024782 5.23736037 4.44583893 -0.09848674 -0.36530670
##      Y30
## -0.80020385

```

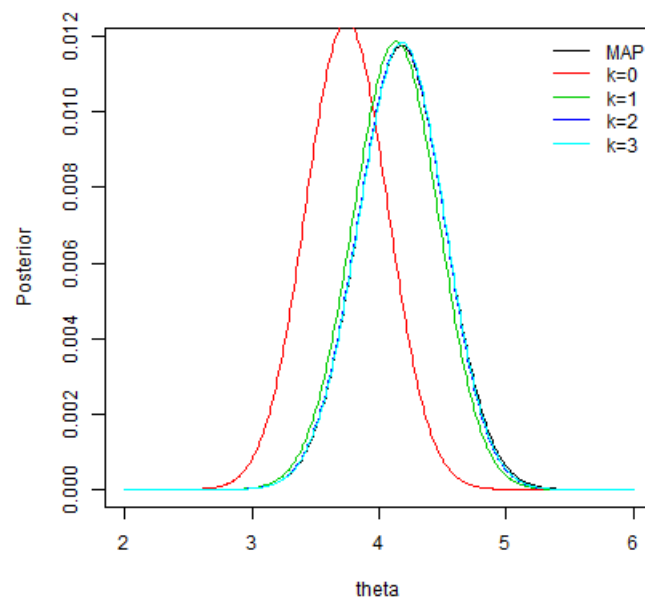
```
##          theta
## theta 0.3389466
```

```
map      <- 4.17919603
```

(13d)

```
posterior <- function(theta,Y,k){
  post <- dnorm(theta,0,sqrt(10^k))
  for(i in 1:length(Y)){
    post<-post*(0.5*dnorm(Y[i],0,1)+
                  0.5*dnorm(Y[i],theta,1))
  }
  return(post/sum(post))}

theta <- seq(2,6,0.01)
map    <- dnorm(theta,map,sd)
plot(theta,map/sum(map),type="l",ylab="Posterior")
lines(theta,posterior(theta,Y,0),col=2)
lines(theta,posterior(theta,Y,1),col=3)
lines(theta,posterior(theta,Y,2),col=4)
lines(theta,posterior(theta,Y,3),col=5)
legend("topright",c("MAP", "k=0", "k=1", "k=2", "k=3"),
      col=1:5,lty=1,bty="n")
```



Even though an improper prior is problematic in this case, the normal prior with large variance performs similar to the Gaussian approximation.

(13e)

```

library(rjags)
data <- list(n=n,Y=Y)

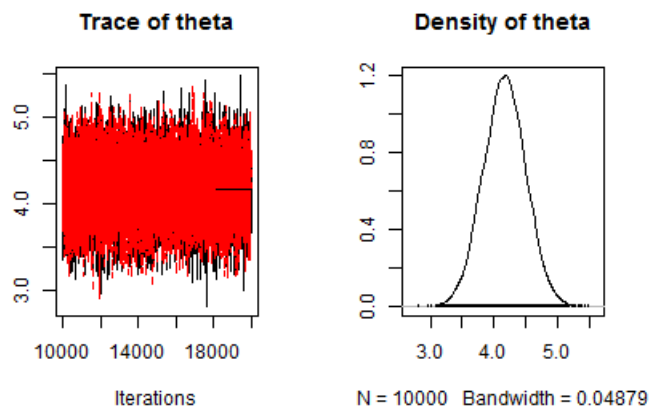
model_string <- textConnection("model{

# Likelihood
for(i in 1:n){
  Y[i] ~ dnorm(B[i]*theta,1)
  B[i] ~ dbern(0.5)
}

# Priors
theta ~ dnorm(0, 0.01)
}")

inits <- list(theta=1)
model <- jags.model(model_string,data = data, inits=inits, n.chains=2,quiet=TRUE)
update(model, 10000, progress.bar="none")
params <- c("theta")
samples <- coda.samples(model,
  variable.names=params,
  n.iter=10000, progress.bar="none")
plot(samples)

```



Convergence looks good and the results are similar to part d.

(15a)

```

library(rjags)
mass <- c(29.9, 1761, 1807, 2984, 3230, 5040, 5654)
age <- c(2, 15, 14, 16, 18, 22, 28)
n <- length(age)
data <- list(mass=mass,age=age,n=n)

model_string <- textConnection("model{

# Likelihood
for(i in 1:n){
  mass[i] ~ dnorm(mu[i],tau)
  mu[i] <- theta[1] + theta[2]*pow(age[i],theta[3])
}

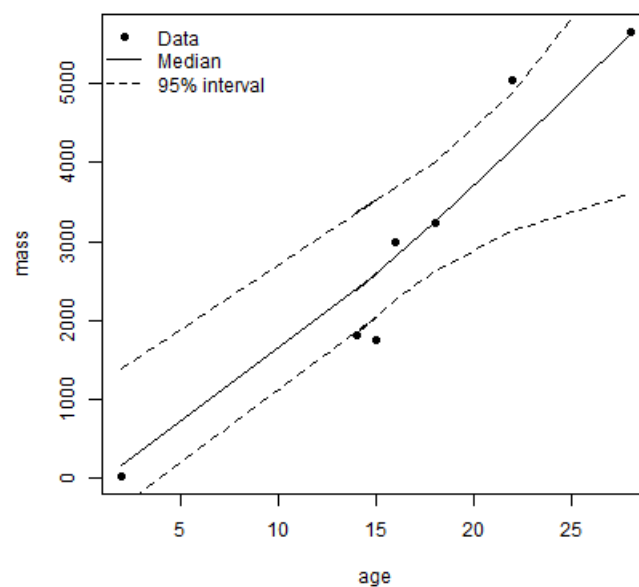
# Priors
theta[1] ~ dnorm(0, 0.00001)
theta[2] ~ dunif(0, 10000)
theta[3] ~ dnorm(0,1)
tau ~ dgamma(0.01,0.01)
}")

model <- jags.model(model_string,data = data, n.chains=2,quiet=TRUE)
update(model, 10000, progress.bar="none")
params <- c("mu","theta")
samples <- coda.samples(model,
  variable.names=params,
  n.iter=10000, progress.bar="none")

q <- summary(samples)$quantiles

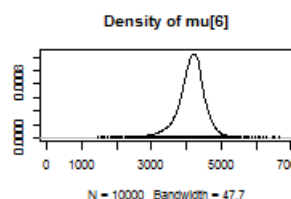
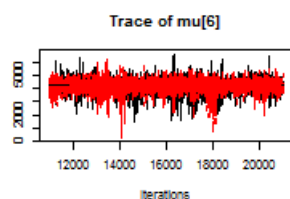
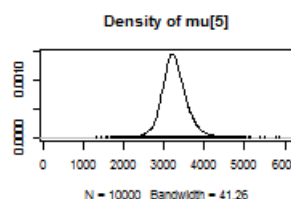
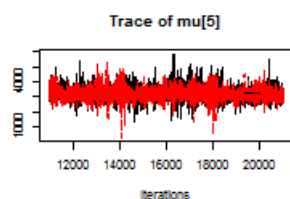
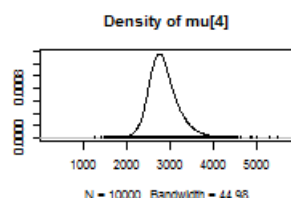
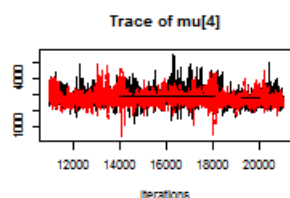
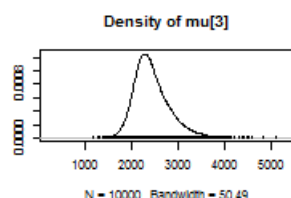
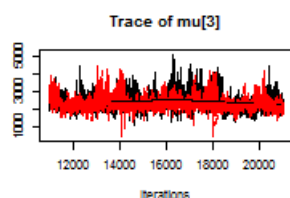
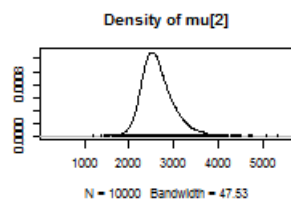
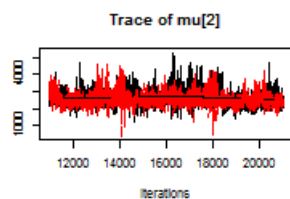
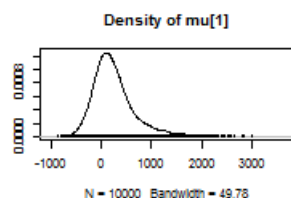
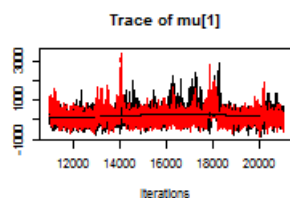
plot(age,mass,pch=19)
lines(age,q[1:n,3]) # Posterior medians
lines(age,q[1:n,1],lty=2) # Posterior 95% interval
lines(age,q[1:n,5],lty=2)
legend("topleft",c("Data","Median","95% interval"),
  lty=c(NA,1,2),pch=c(19,NA,NA),bty="n")

```

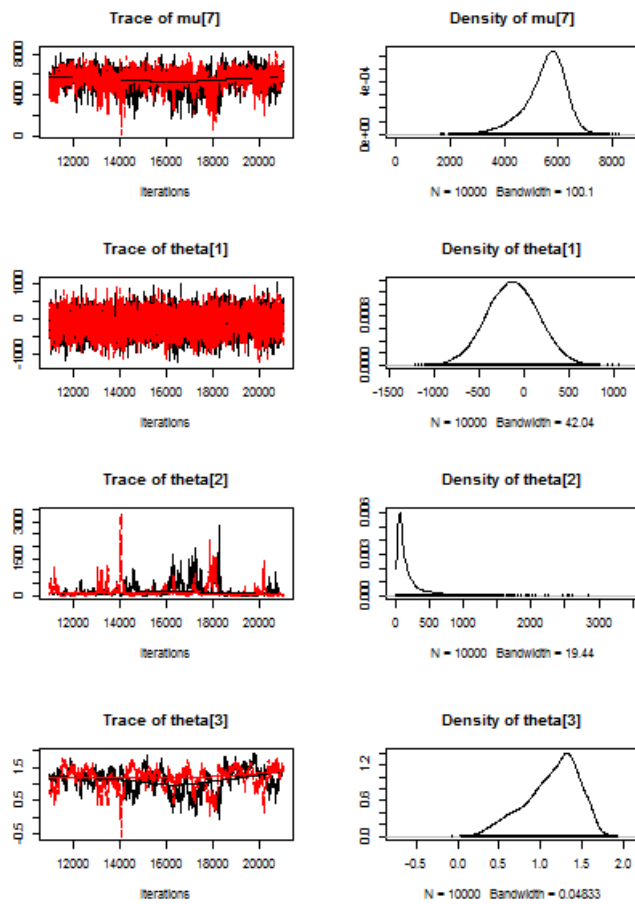


(15b)

```
plot(samples)
```







```
summary(samples)
```

```
##
## Iterations = 11001:21000
## Thinning interval = 1
## Number of chains = 2
## Sample size per chain = 10000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##          Mean      SD Naive SE Time-series SE
## mu[1]      256.082 435.2737 3.077850      33.62855
## mu[2]     2639.986 371.5540 2.627284      30.46857
## mu[3]     2437.227 388.9419 2.750234      36.94359
## mu[4]     2845.571 357.2971 2.526472      24.12871
## mu[5]     3264.930 345.8699 2.445669      10.73452
## mu[6]     4134.656 430.8688 3.046702      13.72068
## mu[7]     5510.568 794.6578 5.619079      70.75311
## theta[1]  -121.491 287.4275 2.032419        4.98096
## theta[2]   215.479 292.6183 2.069124      26.57504
## theta[3]    1.139  0.3305 0.002337        0.04447
##
## 2. Quantiles for each variable:
##
##          2.5%      25%      50%      75%     97.5%
## mu[1]     -363.8349 -23.5302 178.593 432.571 1386.68
## mu[2]     2041.2965 2397.3442 2590.581 2832.836 3521.57
## mu[3]     1829.0203 2176.9567 2377.629 2639.582 3372.65
## mu[4]     2257.2552 2618.8173 2806.125 3030.927 3669.95
## mu[5]     2622.9779 3069.7937 3246.819 3447.838 4012.44
## mu[6]     3138.4372 3942.2632 4173.945 4379.307 4875.10
## mu[7]     3591.0847 5116.0648 5635.422 6032.957 6754.45
## theta[1]  -686.4429 -316.9708 -121.804  73.388  437.62
## theta[2]    26.0520  59.4941 108.357 237.609 1083.00
## theta[3]    0.4038  0.9362  1.196  1.381  1.64
```

```
effectiveSize(samples)
```

```
##      mu[1]      mu[2]      mu[3]      mu[4]      mu[5]      mu[6]
## 172.45780 156.28386 115.72326 234.80936 1095.44094 1043.56345
##      mu[7]  theta[1]  theta[2]  theta[3]
## 125.61708 3515.13968 126.57898  55.19019
```

```
gelman.diag(samples,multivariate=F)
```

```
## Potential scale reduction factors:
##
##          Point est. Upper C.I.
## mu[1]      1.01      1.03
## mu[2]      1.01      1.05
## mu[3]      1.01      1.06
## mu[4]      1.01      1.04
## mu[5]      1.00      1.02
## mu[6]      1.00      1.00
## mu[7]      1.01      1.03
## theta[1]   1.00      1.00
## theta[2]   1.01      1.04
## theta[3]   1.02      1.07
```

Convergence is fine for  $\theta_1$ , but poor for  $\theta_2$  and  $\theta_3$  which have effective sample size less than 100.

(15c) One approach to improving convergence is to simply run the chains longer. This will likely work here because convergence isn't hopeless and the code is fast. A second approach is to simplify the model. The log-linear model

$$\log(\text{mass}) = \beta_0 + \beta_1 \log(\text{age}) + \epsilon$$

is one possibility. Finally,  $\theta_2$  and  $\theta_3$  have posterior correlation -0.88 and so updating them in a block might improve convergence.

(17a) Convergence could be slow because there are more parameters than observations and so it is likely that not all parameters are identifiable.

(17b)

```

library(rjags)
mod <- "model{

  # Likelihood
  for(i in 1:n){
    Y[i] ~ dnorm(mu[i],tau[i])
    mu[i] ~ dnorm(0,theta[1])
    tau[i] ~ dgamma(theta[2],theta[3])
  }

  # Priors
  theta[1] ~ dgamma(eps, eps)
  theta[2] ~ dgamma(eps, eps)
  theta[3] ~ dgamma(eps, eps)
}"

params <- c("theta")

model_string <- textConnection(mod)
data <- list(Y=1:5,n=5,eps=0.1)
model <- jags.model(model_string,data = data, n.chains=2,quiet=TRUE)
update(model, 10000, progress.bar="none")
s <- coda.samples(model,
  variable.names=params,
  n.iter=20000, progress.bar="none")
ESS1 <- effectiveSize(s)

model_string <- textConnection(mod)
data <- list(Y=1:25,n=25,eps=0.1)
model <- jags.model(model_string,data = data, n.chains=2,quiet=TRUE)
update(model, 10000, progress.bar="none")
s <- coda.samples(model,
  variable.names=params,
  n.iter=20000, progress.bar="none")
ESS2 <- effectiveSize(s)

model_string <- textConnection(mod)
data <- list(Y=1:5,n=5,eps=10)
model <- jags.model(model_string,data = data, n.chains=2,quiet=TRUE)
update(model, 10000, progress.bar="none")
s <- coda.samples(model,
  variable.names=params,
  n.iter=20000, progress.bar="none")
ESS3 <- effectiveSize(s)

model_string <- textConnection(mod)
data <- list(Y=1:25,n=25,eps=10)
model <- jags.model(model_string,data = data, n.chains=2,quiet=TRUE)
update(model, 10000, progress.bar="none")
s <- coda.samples(model,
  variable.names=params,
  n.iter=20000, progress.bar="none")
ESS4 <- effectiveSize(s)

ESS1 # n = 5 and eps = 0.1

```

```

## theta[1] theta[2] theta[3]
## 1399.1240 268.1794 625.9386

```

```

ESS2 # n = 25 and eps = 0.1

```

```
## theta[1] theta[2] theta[3]
## 30802.3032 57.2700 195.5023
```

ESS3 #  $n = 5$  and  $eps = 10$

```
## theta[1] theta[2] theta[3]
## 16901.53 12979.49 27441.12
```

ESS4 #  $n = 35$  and  $eps = 10$

```
## theta[1] theta[2] theta[3]
## 10533.82 18625.49 29523.03
```

In the first case with small sample size and uninformative prior convergence isn't great. In the second case the sample size increases, but so does the number of parameters, and so convergence remains problematic. However, increasing the prior information improves convergence for the last two fits.

Processing math: 100%