

# Metropolis + Gibbs sampling for simulated data

## Chapter 3.2.3: Metropolis sampling

Here we use simulated data to verify that the code is working properly. We generate a synthetic dataset and run MCMC to estimate the parameters. Unlike a real data analysis, for a simulation study we know the true values of the parameters and so this provides a check that the code is functioning well.

The data are generated from the logistic regression model

$$Y_i \sim \text{Bernoulli}(p_i) \text{ where } p_i = \frac{1}{1 + \exp(-\sum_{j=1}^p X_{ij}\beta_j)}.$$

The covariates are generated as  $X_{i1} = 1$  for the intercept and  $X_{ij} \sim \text{Normal}(0, 1)$  for  $j > 1$ . The true values of  $\beta_j$  are also sampled from a normal distribution.

We fit the model with prior  $\beta_j \sim \text{Normal}(0, \sigma^2)$  and  $\sigma^2 \sim \text{Gamma}(a, b)$ . The MCMC code using Metropolis steps to update the  $\beta_j$  and a Gibbs step to update  $\sigma^2$ .

## Simulate data

```
n      <- 100
p      <- 20
X      <- cbind(1, matrix(rnorm(n*(p-1)), n, p-1)) # first column for the intercept
true_beta <- rnorm(p, 0, .5)
prob   <- 1/(1+exp(-X%*%true_beta))
Y      <- rbinom(n, 1, prob)
```

## Initialize

```
# Create matrix to store the samples

S      <- 25000
samples <- matrix(NA, S, p+1)

# Initial values

beta   <- rep(0, p)
sigma  <- 1

# priors:

a      <- 0.1
b      <- 0.1

# candidate standard deviation:

can_sd <- 0.1
```

## Define the log posterior as a function

```
log_post_beta <- function(Y, X, beta, sigma){
  prob <- 1/(1+exp(-X%*%beta))
  like <- sum(dbinom(Y, 1, prob, log=TRUE))
  prior <- dnorm(beta[1], 0, 10, log=TRUE) + # Intercept
            sum(dnorm(beta[-1], 0, sigma, log=TRUE)) # Slopes
  return(like+prior)}
```

## Metropolis sampling

```

for(s in 1:S){

  # Metropolis for beta
  for(j in 1:p){
    can  <- beta
    can[j] <- rnorm(1,beta[j],can_sd)
    logR  <- log_post_beta(Y,X,can,sigma)-
              log_post_beta(Y,X,beta,sigma)
    if(log(runif(1))<logR){
      beta <- can
    }
  }

  # Gibbs for sigma
  sigma <- 1/sqrt(rgamma(1,(p-1)/2+a,sum(beta[-1]^2)/2+b))

  samples[s,] <- c(beta,sigma)
}

```

## Plot the results

The boxplots are the posterior distributions for the  $\beta_j$ , and the points are the true values used to generate the data. The posteriors include the truth for most of the parameters indicating the alorithm is working well.

```

boxplot(samples[,1:p],outline=FALSE,xlab="Index, j",ylab=expression(beta[j]))
points(true_beta,pch=19,cex=1.25)

```

