

Comparison of Bayesian software

In this example, we compare JAGS to other Bayesian software to samples from a random slopes models. Let Y_{ij} be the j^{th} measurement of jaw bone density for patient i . In this model we allow bone density to increase linearly in time and each patient has their own slope and intercept. The random slope model is

$$Y_{ij} | \alpha_{i1}, \alpha_{i2} \sim \text{Normal}(\alpha_{i1} + \text{age}_j \alpha_{i2}, \sigma_3^2) \text{ where } (\alpha_{i1}, \alpha_{i2})^T \sim \text{Normal}(\mu, \Sigma).$$

The random effects α_{i1} and α_{i2} are the subject-specific intercept and slope, respectively. The population of intercepts and slopes are assumed to be $\alpha_{ji} \sim \text{Normal}(\mu_j, \sigma_j)$.

This model is fit below using


1. JAGS
2. OpenBUGS
3. STAN
4. NIMBLE

Load jaw data

```
m <- 4
n <- 20
age <- c(8.0, 8.5, 9.0, 9.5)
Y <- c(47.8, 48.8, 49.0, 49.7,
      46.4, 47.3, 47.7, 48.4,
      46.3, 46.8, 47.8, 48.5,
      45.1, 45.3, 46.1, 47.2,
      47.6, 48.5, 48.9, 49.3,
      52.5, 53.2, 53.3, 53.7,
      51.2, 53.0, 54.3, 54.5,
      49.8, 50.0, 50.3, 52.7,
      48.1, 50.8, 52.3, 54.4,
      45.0, 47.0, 47.3, 48.3,
      51.2, 51.4, 51.6, 51.9,
      48.5, 49.2, 53.0, 55.5,
      52.1, 52.8, 53.7, 55.0,
      48.2, 48.9, 49.3, 49.8,
      49.6, 50.4, 51.2, 51.8,
      50.7, 51.7, 52.7, 53.3,
      47.2, 47.7, 48.4, 49.5,
      53.3, 54.6, 55.1, 55.3,
      46.2, 47.5, 48.1, 48.4,
      46.3, 47.6, 51.3, 51.8)

Y <- matrix(Y,20,4,byrow=TRUE)

plot(NA,xlim=range(age),ylim=range(Y),xlab="Age",ylab="Bone density")
for(i in 1:n){
  lines(age,Y[i,])
  points(age,Y[i,],pch=19)
}
```



(1) Mixed model in JAGS

```
library(rjags)
```

```
## Loading required package: coda
```

```
## Linked to JAGS 4.2.0
```

```
## Loaded modules: basemod,bugs
```

```

set.seed(0820)

tick <- proc.time()[3]
model_string <- textConnection("model{
  # Likelihood
  for(i in 1:n){for(j in 1:m){
    Y[i,j] ~ dnorm(alpha1[i]+alpha2[i]*age[j],tau3)
  }}

  # Random effects
  for(i in 1:n){
    alpha1[i] ~ dnorm(mu1,tau1)
    alpha2[i] ~ dnorm(mu2,tau2)
  }

  mu1 ~ dnorm(0,0.0001)
  mu2 ~ dnorm(0,0.0001)
  tau1 ~ dgamma(0.1,0.1)
  tau2 ~ dgamma(0.1,0.1)
  tau3 ~ dgamma(0.1,0.1)
}")

data <- list(Y=Y,age=age,n=n,m=m)
params <- c("mu1","mu2","tau1","tau2","tau3")
model <- jags.model(model_string,data = data, n.chains=2,quiet=TRUE)
update(model, 10000, progress.bar="none")
samples <- coda.samples(model, variable.names=params,
  n.iter=90000, progress.bar="none")

tock <- proc.time()[3]
tock-tick

```

```

## elapsed
##      1.83

```

```
effectiveSize(samples)
```

```

##      mu1      mu2      tau1      tau2      tau3
## 292.9739 335.4025 3406.6998 6325.8104 15608.3856

```

(2) Mixed model in OpenBUGS

```
library(R2OpenBUGS)
```

```
## Warning: package 'R2OpenBUGS' was built under R version 3.3.3
```

```

set.seed(0820)

tick <- proc.time()[3]

model_string <- function(){
  # Likelihood
  for(i in 1:n){for(j in 1:m){
    Y[i,j] ~ dnorm(mn[i,j],tau3)
    mn[i,j] <- alpha1[i]+alpha2[i]*age[j]
  }}

  # Random effects
  for(i in 1:n){
    alpha1[i] ~ dnorm(mu1,tau1)
    alpha2[i] ~ dnorm(mu2,tau2)
  }

  # Priors
  mu1 ~ dnorm(0,0.0001)
  mu2 ~ dnorm(0,0.0001)
  tau1 ~ dgamma(0.1,0.1)
  tau2 ~ dgamma(0.1,0.1)
  tau3 ~ dgamma(0.1,0.1)
}

data <- list(Y=Y,age=age,n=n,m=m)
params <- c("mu1","mu2","tau1","tau2","tau3")
inits <- function(){list(mu1=0,mu2=0,tau1=.1,tau2=.2,tau3=.2)}
fit <- bugs(model.file=model_string,
            data=data,inits=inits,
            parameters.to.save=params,DIC=FALSE,
            n.iter=100000,n.chains=2,n.burnin=10000)

tock <- proc.time()[3]
tock-tick

```

```

## elapsed
## 10.17

```

```
fit$summary[,9]
```

```

##      mu1      mu2    tau1    tau2    tau3
## 1300    1200 180000  11000 180000

```

(3) Mixed model in STAN

```
library(rstan)
```

```
## Warning: package 'rstan' was built under R version 3.3.3
```

```
## Loading required package: ggplot2
```

```
## Use suppressPackageStartupMessages() to eliminate package startup
## messages.
```

```
## Loading required package: StanHeaders
```

```
## Warning: package 'StanHeaders' was built under R version 3.3.3
```

```
## rstan (Version 2.17.3, GitRev: 2e1f913d3ca3)
```

```
## For execution on a local, multicore CPU with excess RAM we recommend calling
## options(mc.cores = parallel::detectCores()).
## To avoid recompilation of unchanged Stan programs, we recommend calling
## rstan_options(auto_write = TRUE)
```

```
##
## Attaching package: 'rstan'
```

```
## The following object is masked from 'package:R2OpenBUGS':
##
##     monitor
```

```
## The following object is masked from 'package:coda':
##
##     traceplot
```

```
set.seed(0820)

tick <- proc.time()[3]
stan_model <- "

  data {
    int<lower=0> n;
    int<lower=0> m;
    vector [m] age;
    matrix [n,m] Y;
  }

  parameters {
    vector [n] alpha1;
    vector [n] alpha2;
    real mu1;
    real mu2;
    real<lower=0> sigma3;
    real<lower=0> sigma2;
    real<lower=0> sigma1;
  }

  model {
    real mu;
    alpha1 ~ normal(0,sigma1);
    alpha2 ~ normal(0,sigma2);
    sigma1 ~ cauchy(0.0,1000);
    sigma2 ~ cauchy(0.0,1000);
    sigma3 ~ cauchy(0.0,1000);
    mu1 ~ normal(0,10000);
    mu2 ~ normal(0,10000);

    for(i in 1:n){for(j in 1:m){
      mu = alpha1[i] + alpha2[i]*age[j];
      Y[i,j] ~ normal(mu,sigma3);
    }}
  }
"

data <- list(Y=Y,age=age,n=n,m=m)
fit_stan <- stan(model_code = stan_model,
  data = data,
  chains=2, warmup = 10000, iter = 100000)
```

```
## In file included from C:/Program Files/R/R-3.3.2/library/BH/include/boost/config.hpp:39:0,
##      from C:/Program Files/R/R-3.3.2/library/BH/include/boost/math/tools/config.hpp:13,
##      from C:/Program Files/R/R-3.3.2/library/StanHeaders/include/stan/math/rev/core/var.hpp:7,
##      from C:/Program Files/R/R-3.3.2/library/StanHeaders/include/stan/math/rev/core/gevv_vvv_var
```

```

##          from C:/Program Files/R/R-3.3.2/library/StanHeaders/include/stan/math/rev/core.hpp:12,
##          from C:/Program Files/R/R-3.3.2/library/StanHeaders/include/stan/math/rev/mat.hpp:4,
##          from C:/Program Files/R/R-3.3.2/library/StanHeaders/include/stan/math.hpp:4,
##          from C:/Program Files/R/R-3.3.2/library/StanHeaders/include/src/stan/model/model_header.hpp:8,
##          from file175c253b4319.cpp:8:
## C:/Program Files/R/R-3.3.2/library/BH/include/boost/config/compiler/gcc.hpp:186:0: warning: "BOOST_NO_CXX11_
## # define BOOST_NO_CXX11_RVALUE_REFERENCES
## ^
## <command-line>:0:0: note: this is the location of the previous definition
## In file included from C:/Program Files/R/R-3.3.2/library/StanHeaders/include/stan/math/rev/core.hpp:44:0,
##          from C:/Program Files/R/R-3.3.2/library/StanHeaders/include/stan/math/rev/mat.hpp:4,
##          from C:/Program Files/R/R-3.3.2/library/StanHeaders/include/stan/math.hpp:4,
##          from C:/Program Files/R/R-3.3.2/library/StanHeaders/include/src/stan/model/model_header.hpp:8,
##          from file175c253b4319.cpp:8:
## C:/Program Files/R/R-3.3.2/library/StanHeaders/include/stan/math/rev/core/set_zero_all_adjoints.hpp:14:17: v
##     static void set_zero_all_adjoints() {
##         ^
##
## SAMPLING FOR MODEL '4bf61f87d83553920d54d39a64269860' NOW (CHAIN 1).
##
## Gradient evaluation took 0 seconds
## 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Adjust your expectations accordingly!
##
##
## Iteration:    1 / 100000 [ 0%] (Warmup)
## Iteration: 10000 / 100000 [ 10%] (Warmup)
## Iteration: 10001 / 100000 [ 10%] (Sampling)
## Iteration: 20000 / 100000 [ 20%] (Sampling)
## Iteration: 30000 / 100000 [ 30%] (Sampling)
## Iteration: 40000 / 100000 [ 40%] (Sampling)
## Iteration: 50000 / 100000 [ 50%] (Sampling)
## Iteration: 60000 / 100000 [ 60%] (Sampling)
## Iteration: 70000 / 100000 [ 70%] (Sampling)
## Iteration: 80000 / 100000 [ 80%] (Sampling)
## Iteration: 90000 / 100000 [ 90%] (Sampling)
## Iteration: 100000 / 100000 [100%] (Sampling)
##
## Elapsed Time: 15.162 seconds (Warm-up)
##              189.403 seconds (Sampling)
##              204.565 seconds (Total)
##
## SAMPLING FOR MODEL '4bf61f87d83553920d54d39a64269860' NOW (CHAIN 2).
##
## Gradient evaluation took 0 seconds
## 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Adjust your expectations accordingly!
##
##
## Iteration:    1 / 100000 [ 0%] (Warmup)
## Iteration: 10000 / 100000 [ 10%] (Warmup)
## Iteration: 10001 / 100000 [ 10%] (Sampling)
## Iteration: 20000 / 100000 [ 20%] (Sampling)
## Iteration: 30000 / 100000 [ 30%] (Sampling)
## Iteration: 40000 / 100000 [ 40%] (Sampling)
## Iteration: 50000 / 100000 [ 50%] (Sampling)
## Iteration: 60000 / 100000 [ 60%] (Sampling)
## Iteration: 70000 / 100000 [ 70%] (Sampling)
## Iteration: 80000 / 100000 [ 80%] (Sampling)
## Iteration: 90000 / 100000 [ 90%] (Sampling)
## Iteration: 100000 / 100000 [100%] (Sampling)
##
## Elapsed Time: 16.439 seconds (Warm-up)
##              173.239 seconds (Sampling)

```

```
## 189.678 seconds (Total)
```

```
tock <- proc.time()[3]  
tock-tick
```

```
## elapsed  
## 424.27
```

```
summary(fit_stan)$summary[41:45,9:10]
```

```
##          n_eff      Rhat  
## mu1    180000.00 1.0000077  
## mu2    180000.00 1.0000082  
## sigma3  79991.02 0.9999959  
## sigma2 180000.00 0.9999899  
## sigma1 180000.00 0.9999975
```

(4) Mixed model in NIMBLE

```
library(nimble)
```

```
## nimble version 0.6-11 is loaded.  
## For more information on NIMBLE and a User Manual,  
## please visit http://R-nimble.org.
```

```
##  
## Attaching package: 'nimble'
```

```
## The following object is masked from 'package:stats':  
##  
## simulate
```

```

set.seed(0820)

tick <- proc.time()[3]

model_string <- nimbleCode({
  # Likelihood
  for(i in 1:n){for(j in 1:m){
    Y[i,j] ~ dnorm(mn[i,j],tau3)
    mn[i,j] <- alpha1[i]+alpha2[i]*age[j]
  }}

  # Random effects
  for(i in 1:n){
    alpha1[i] ~ dnorm(mu1,tau1)
    alpha2[i] ~ dnorm(mu2,tau2)
  }

  # Priors
  mu1 ~ dnorm(0,0.0001)
  mu2 ~ dnorm(0,0.0001)
  tau1 ~ dgamma(0.1,0.1)
  tau2 ~ dgamma(0.1,0.1)
  tau3 ~ dgamma(0.1,0.1)
})

consts <- list(n=n,m=m,age=age)
data <- list(Y=Y)
inits <- function(){list(mu1=rnorm(1),mu2=rnorm(1),tau1=10,tau2=10,tau3=10)}
samples <- nimbleMCMC(model_string, data = data, inits = inits,
                      constants=consts,
                      monitors = c("mu1", "mu2", "tau1", "tau2", "tau3"),
                      samplesAsCodaMCMC=TRUE,WAIC=FALSE,
                      niter = 100000, nburnin = 10000, nchains = 2)

```

```
## defining model...
```

```
## building model...
```

```
## setting data and initial values...
```

```
## running calculate on model (any error reports that follow may simply
## reflect missing values in model variables) ...
```

```
##
```

```
## checking model sizes and dimensions...
```

```
## This model is not fully initialized. This is not an error. To see which
## variables are not initialized, use model$initializeInfo(). For more
## information on model initialization, see help(modelInitialization).
```

```
##
```

```
## checking model calculations...
```

```
## NAs were detected in model variables: alpha1, logProb_alpha1, alpha2, logProb_alpha2, mn, logProb_Y.
```

```
## model building finished.
```

```
## compiling... this may take a minute. Use 'showCompilerOutput = TRUE' to see C++ compiler details.
```



```
## compilation finished.
```

```
## running chain 1...
```

```
## |-----|-----|-----|-----|
## |-----|-----|-----|-----|
```

```
## running chain 2...
```

```
## |-----|-----|-----|-----|
## |-----|-----|-----|-----|
```

```
tock <- proc.time()[3]
tock-tick
```

```
## elapsed
## 26.49
```

```
effectiveSize(samples)
```

```
##      mu1      mu2      tau1      tau2      tau3
## 283.4198 310.7662 3448.1076 6237.2704 13035.0535
```