

Analysis of tyrannosaurid growth curves

Chapter 7: Case studies using hierarchical modelling

We analyze the data from 20 fossils to estimate the growth curves of four tyrannosaurid species: *Albertosaurus*, *Daspletosaurus*, *Gorgosaurus* and *Tyrannosaurus*. The data are taken from Table 1 of

Erickson, GM et al (2004). Gigantism and comparative life-history parameters of tyrannosaurid dinosaurs. *Nature*, **430**, 772-775.

Let Y_{ij} and X_{ij} be the body mass and age, respectively, of sample i from species $j = 1, \dots, 4$. We model the data as $Y_{ij} = f_j(X_{ij})\epsilon_{ij}$ where f_j is the true growth curve for species j and ϵ_{ij} is multiplicative error. We use multiplicative error rather than additive error because variation in the population likely increases with mass/age. Assuming the errors are log-normal the model becomes $\log(Y_{ij}) \sim \text{Normal}(\log[f_j(X_{ij})] - \sigma_{2j}^2/2, \sigma_{2j}^2)$, where σ_{2j}^2 is the error variance for species j and is included in the mean so that the multiplicative errors ϵ_{ij} have mean one. In the analysis below we compare four models that vary by the functional form of the growth curves f_j (either log-linear or logistic) and how informative about parameters is shared across species (either not pooled or pooled via Gaussian random effects).

Load and plot the data

```
dat <- read.csv("Erickson.csv")
dat[15:16,] <- dat[16:15,] #order by age
dat
```

##	Taxon	Spec...	Age	Mass
## 1	Gorgosaurus	FMNH PR2211	5	127.0
## 2	Gorgosaurus	RTMP 86.144.1	7	229.0
## 3	Gorgosaurus	RTMP 99.33.1	14	607.0
## 4	Gorgosaurus	RTMP 73.30.1	14	747.0
## 5	Gorgosaurus	RTMP 94.12.602	18	1105.0
## 6	Albertosaurus	RTMP 2002.45.46	2	50.3
## 7	Albertosaurus	RTMP 86.64.01	15	762.0
## 8	Albertosaurus	USNM 12814/AMNH 5428	18	1013.0
## 9	Albertosaurus	AMNH 5432	22	1282.0
## 10	Albertosaurus	RTMP 81.10.1	24	1142.0
## 11	Daspletosaurus	RTMP 94.143.1	10	496.0
## 12	Daspletosaurus	AMNH 5438	17	1518.0
## 13	Daspletosaurus	FMNH PR 308	21	1791.0
## 14	Tyrannosaurus	LACM 28471	2	29.9
## 15	Tyrannosaurus	LACM 23845	14	1807.0
## 16	Tyrannosaurus	AMNH 30564	15	1761.0
## 17	Tyrannosaurus	ICM 2001.90.1	16	2984.0
## 18	Tyrannosaurus	RTMP 81.6.1	18	3230.0
## 19	Tyrannosaurus	RTMP 81.12.1	22	5040.0
## 20	Tyrannosaurus	FMNH PR 2081	28	5654.0

```

taxon      <- dat[,1]
age        <- dat[,3]
mass       <- dat[,4]

# Plot the data


plot(NA,xlim=range(age),ylim=range(mass),
     main="", xlab="Age (years)", ylab="Body Mass (kg)")
points(age[taxon=="Albertosaurus"], mass[taxon=="Albertosaurus"], pch=1,cex=1.5)
points(age[taxon=="Daspletosaurus"],mass[taxon=="Daspletosaurus"],pch=2,cex=1.5)
points(age[taxon=="Gorgosaurus"],   mass[taxon=="Gorgosaurus"],   pch=3,cex=1.5)
points(age[taxon=="Tyrannosaurus"], mass[taxon=="Tyrannosaurus"], pch=4,cex=1.5)

# Fitted growth curves (from the original paper)

x <- seq(0,30,.1)
lines(x,1218/(1+exp(-0.43*(x-14.1)))+5, lty=1)
lines(x,1728/(1+exp(-0.44*(x-12.1)))+5, lty=2)
lines(x,1234/(1+exp(-0.38*(x-12.4)))+5, lty=3)
lines(x,5551/(1+exp(-0.57*(x-16.1)))+5, lty=4)

legend("topleft",c("Albertosaurus","Daspletosaurus","Gorgosaurus","Tyrannosaurus"),
      lty=1:4, pch=1:4,bty="n",cex=1.5)

```



Put the data in JAGS format

```

library(rjags)
set.seed(0820)

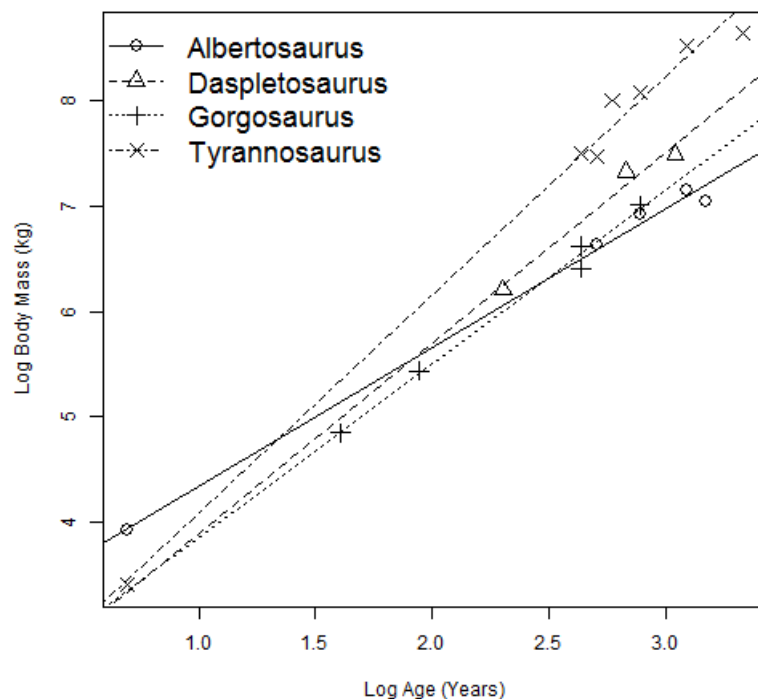
y      <- log(mass)
x      <- log(age)
n      <- length(y)
sp     <- as.numeric(taxon)
names  <- c("Albertosaurus","Daspletosaurus","Gorgosaurus","Tyrannosaurus")
data   <- list(y=y,x=x,sp=sp,n=n,N=4)
data

```

```
## $y
## [1] 4.844187 5.433722 6.408529 6.616065 7.007601 3.918005 6.635947
## [8] 6.920672 7.156177 7.040536 6.206576 7.325149 7.490529 3.397858
## [15] 7.499423 7.473637 8.001020 8.080237 8.525161 8.640119
##
## $x
## [1] 1.6094379 1.9459101 2.6390573 2.6390573 2.8903718 0.6931472 2.7080502
## [8] 2.8903718 3.0910425 3.1780538 2.3025851 2.8332133 3.0445224 0.6931472
## [15] 2.6390573 2.7080502 2.7725887 2.8903718 3.0910425 3.3322045
##
## $sp
## [1] 3 3 3 3 3 1 1 1 1 1 2 2 2 4 4 4 4 4 4 4
##
## $n
## [1] 20
##
## $N
## [1] 4
```

```
burn    <- 10000
iters   <- 100000
thin    <- 10

plot(x,y,pch=sp,cex=1.5,xlab="Log Age (Years)",ylab="Log Body Mass (kg)")
for(j in 1:4){
  b <- lm(y[sp==j]~x[sp==j])
  abline(b[1],b[2],lty=j)
}
legend("topleft",names,lty=1:4, pch=1:4,bty="n",cex=1.5)
```



Summary: The log of age versus the log of mass is roughly linear for all four taxa.

Model 1: Unpooled analysis with log-linear growth curves

The exponential growth model is $f_j(\text{age}) = A_j \text{age}^{b_j}$ so that $\log[f_j(\text{age})] = a_j + b_j \log(\text{age})$, where $\log(A_j) = a_j$. This is clearly unrealistic when applied to the entire life-span of a species because body mass should plateau, but may be reasonable to model growth early in life. In this analysis the priors are $a_j, b_j \sim \text{Normal}(0, 10)$ and $\sigma_{2j}^2 \sim \text{InvGamma}(0.1, 0.1)$ to give a separate regression for each species.

```

model_string1 <- textConnection("model{
  for(i in 1:n){
    y[i] ~ dnorm(muY[i],tau[sp[i]])
    muY[i] <- a[sp[i]] + b[sp[i]]*x[i] - 0.5/tau[sp[i]]
  }

  for(j in 1:N){
    a[j] ~ dnorm(0.0,0.1)
    b[j] ~ dnorm(0.0,0.1)
    tau[j] ~ dgamma(0.1,0.1)
  }

  for(age in 1:30){for(j in 1:4){
    fitted[age,j] <- exp(a[j] + b[j]*log(age))
  }}
}")

model1 <- jags.model(model_string1,data = data, quiet=TRUE,n.chains=2)
update(model1, burn, progress.bar="none")
samples1 <- coda.samples(model1, variable.names=c("a","b","fitted","tau"), n.iter=iters, thin=thin, progress.b

ESS <- effectiveSize(samples1)
ESS[which.min(ESS)]

```

```

##      a[2]
## 3922.673

```

```
ESS[which.max(ESS)]
```

```

## fitted[9,3]
##      20000

```

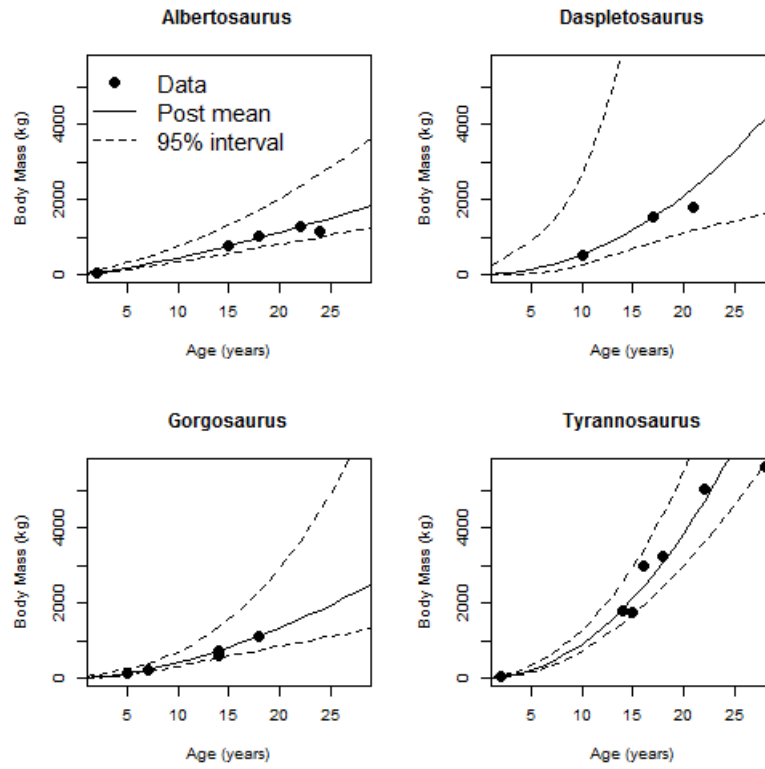
```

fit1 <- summary(samples1)$quantiles[1:120+8,3]
lo1 <- summary(samples1)$quantiles[1:120+8,1]
hi1 <- summary(samples1)$quantiles[1:120+8,5]
id <- rep(1:4,each=30)

par(mfrow=c(2,2))
for(j in 1:4){
  plot(NA,xlim=range(age),ylim=range(mass),
       xlab="Age (years)", ylab="Body Mass (kg)",
       main=names[j])
  points(age[sp==j],mass[sp==j],pch=19,cex=1.5)
  lines(1:30,fit1[id==j],lty=1)
  lines(1:30,lo1[id==j],lty=2)
  lines(1:30,hi1[id==j],lty=2)

  if(j==1){
    legend("topLeft",c("Data", "Post mean", "95% interval"),
          pch=c(19,NA,NA),lty=c(NA,1,2),cex=1.5,bty="n")
  }
}

```



Summary: Convergence is OK (large ESS) and the fit to the data is decent. However, the uncertainty in the mean growth curves is very large.

Model 2: Pooled analysis with log-linear growth curves

This model is the same as Model 1 except that all species have the same variance, $\sigma_{2j} = \sigma^2$, and the regression coefficients are modeled hierarchically. For example, for intercepts we have $a_j \sim \text{Normal}(\mu_a, \sigma_{2a}^2)$, where $\mu_a \sim \text{Normal}(0, 10)$ and $\sigma_{2a}^2 \sim \text{InvGamma}(0.1, 0.1)$. The slopes b_j are modelled similarly.

```

model_string2 <- textConnection("model{
  for(i in 1:n){
    y[i] ~ dnorm(muY[i],tau[1])
    muY[i] <- a[sp[i]] + b[sp[i]]*x[i] - 0.5/tau[1]
  }

  for(j in 1:N){
    a[j] ~ dnorm(mu[1],tau[2])
    b[j] ~ dnorm(mu[2],tau[3])
  }

  for(k in 1:2){
    mu[k] ~ dnorm(0.0,0.1)
  }
  for(k in 1:3){
    tau[k] ~ dgamma(0.1,0.1)
  }

  for(age in 1:30){for(j in 1:4){
    fitted[age,j] <- exp(a[j] + b[j]*log(age))
  }}
}")

model2 <- jags.model(model_string2,data = data,quiet=TRUE, n.chains=2)
update(model2, burn, progress.bar="none")
samples2 <- coda.samples(model2, variable.names=c("a","b","fitted","tau"), n.iter=iters, thin=thin, progress.b

ESS <- effectiveSize(samples2)
ESS[which.min(ESS)]

```

```

##      b[2]
## 4140.49

```

```
ESS[which.max(ESS)]
```

```

## fitted[17,4]
##      20723.92

```

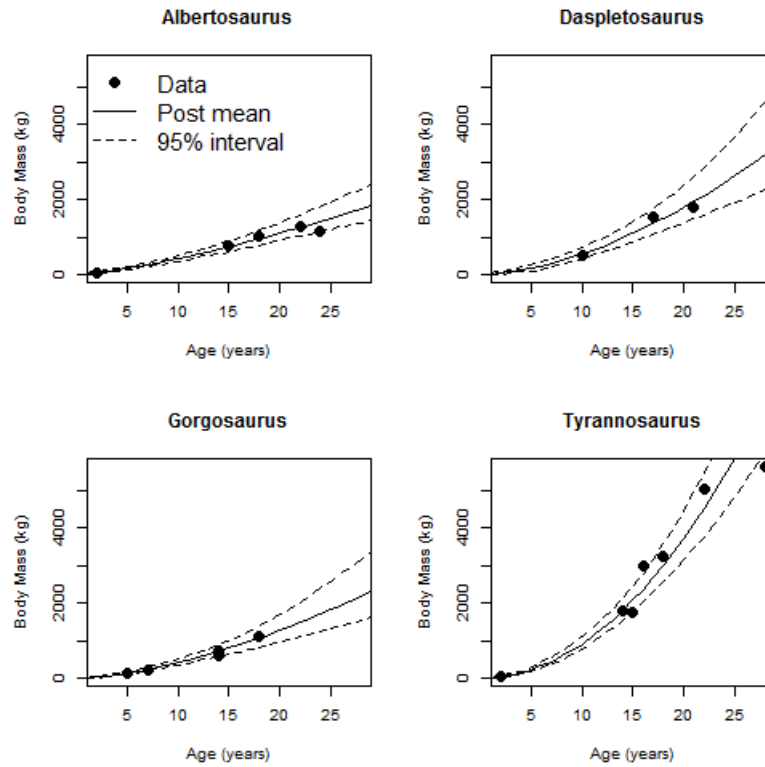
```

fit2 <- summary(samples2)$quantiles[1:120+8,3]
lo2 <- summary(samples2)$quantiles[1:120+8,1]
hi2 <- summary(samples2)$quantiles[1:120+8,5]
id <- rep(1:4,each=30)

par(mfrow=c(2,2))
for(j in 1:4){
  plot(NA,xlim=range(age),ylim=range(mass),
       xlab="Age (years)", ylab="Body Mass (kg)",
       main=names[j])
  points(age[sp==j],mass[sp==j],pch=19,cex=1.5)
  lines(1:30,fit2[id==j],lty=1)
  lines(1:30,lo2[id==j],lty=2)
  lines(1:30,hi2[id==j],lty=2)

  if(j==1){
    legend("topleft",c("Data", "Post mean", "95% interval"),
          pch=c(19,NA,NA),lty=c(NA,1,2),cex=1.5,bty="n")
  }
}

```



Summary: The posterior mean growth curve is similar to Model 1, but the posterior variance is small; this is the effect of pooling information across curves.

Model 3: Non-pooled analysis with logistic growth curves

The non-linear model is $(x = \log(\text{age})) f_j(x) = a_j + b_j \frac{1}{1 + \exp[-(x - c_j)/d_j]}$, which is an increasing function that plateaus as age increases. The growth curve for species j determined by four parameters:

1. a_j is the average mass at age 0
2. b_j is the lifetime expected gain in mass
3. $\log(c_j)$ is the age at which the species is half way to maturity
4. d_j determines the rate of increase

For the curve to be positive and increasing for all ages, we must have $a_j > 0$, $b_j > a_j$ and $d_j > 0$. We enforce these constraints by representing the growth-curve parameters in terms of unconstrained parameters $\alpha_{j1}, \dots, \alpha_{j4}$ as:

1. $a_j = \exp(\alpha_{j1})$
2. $b_j = \exp(\alpha_{j2})$
3. $c_j = \alpha_{j3}$
4. $d_j = \exp(\alpha_{j4})$

In this analysis we use uninformative prior for all parameters, separately by species: $\alpha_{jk} \sim \text{Normal}(0, 10)$ and $\sigma_{2j} \sim \text{InvGamma}(0.1, 0.1)$.

Note that for any value of the α_{jk} , the growth curve is positive and increasing for all ages.

```

model_string3 <- textConnection("model{
  for(i in 1:n){
    y[i]      ~ dnorm(muY[i],tau[sp[i]])
    muY[i]    <- log(a[sp[i]] + b[sp[i]]/(1+exp(-part[i]))) - 0.5/tau[sp[i]]
    part[i]   <- (x[i]-c[sp[i]])/d[sp[i]]
  }

  for(j in 1:N){
    a[j]      <- exp(alpha[j,1])
    b[j]      <- exp(alpha[j,2])
    c[j]      <- alpha[j,3]
    d[j]      <- exp(alpha[j,4])

    for(k in 1:4){alpha[j,k] ~ dnorm(0.0,0.1)}
    tau[j] ~ dgamma(0.1,0.1)
  }

  for(age in 1:30){for(j in 1:4){
    PART[age,j] <- (log(age)-c[j])/d[j]
    fitted[age,j] <- a[j] + b[j]/(1+exp(-PART[age,j]))
  }}

}")

model3 <- jags.model(model_string3,data = data,quiet=TRUE, n.chains=2)
update(model3, burn, progress.bar="none")

samples3 <- coda.samples(model3, variable.names=c("a","b","c","d","fitted","tau"), n.iter=iters, thin=thin, pr

ESS <- effectiveSize(samples3)
ESS[which.min(ESS)]

```

```

##      c[4]
## 1232.332

```

```
ESS[which.max(ESS)]
```

```

## fitted[14,3]
##      20000

```

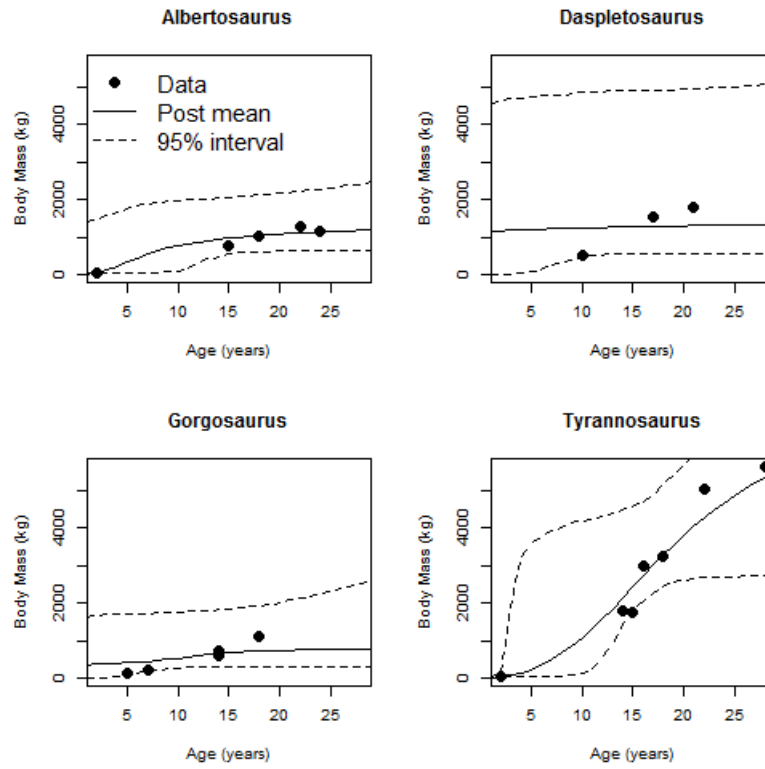
```

fit3 <- summary(samples3)$quantiles[1:120+16,3]
lo3 <- summary(samples3)$quantiles[1:120+16,1]
hi3 <- summary(samples3)$quantiles[1:120+16,5]
id <- rep(1:4,each=30)

par(mfrow=c(2,2))
for(j in 1:4){
  plot(NA,xlim=range(age),ylim=range(mass),
       xlab="Age (years)", ylab="Body Mass (kg)",
       main=names[j])
  points(age[sp==j],mass[sp==j],pch=19,cex=1.5)
  lines(1:30,fit3[id==j],lty=1)
  lines(1:30,lo3[id==j],lty=2)
  lines(1:30,hi3[id==j],lty=2)

  if(j==1){
    legend("topleft",c("Data", "Post mean", "95% interval"),
          pch=c(19,NA,NA),lty=c(NA,1,2),cex=1.5,bty="n")
  }
}

```

Summary: Convergence is questionable (low ESS) and the 95% intervals for the curves are very wide. It appears that there is not enough data to estimate a non-linear growth curve separately by species.

Model 4: Pooled analysis of logistic growth curves

This model is the same as Model 3 except that all species have the same variance, $\sigma_{2j} = \sigma^2$, and the regression coefficients are modeled hierarchically. The random effect distributions are $\log(a_{jk}) \sim \text{Normal}(\mu_k, \sigma_{2k}^2)$, where $\mu_k \sim \text{Normal}(0, 10)$ and $\sigma_{2k}^2 \sim \text{InvGamma}(0.1, 0.1)$.

```

model_string4 <- textConnection("model{
  for(i in 1:n){
    y[i] ~ dnorm(muY[i],tau[5])
    muY[i] <- log(a[sp[i]] + b[sp[i]]/(1+exp(-part[i]))) - 0.5/tau[5]
    part[i] <- (x[i]-c[sp[i]])/d[sp[i]]
  }

  for(j in 1:N){
    a[j] <- exp(alpha[j,1])
    b[j] <- exp(alpha[j,2])
    c[j] <- alpha[j,3]
    d[j] <- exp(alpha[j,4])

    for(k in 1:4){alpha[j,k] ~ dnorm(mu[k],tau[k])}
  }

  for(j in 1:4){
    mu[j] ~ dnorm(0,0.1)
    tau[j] ~ dgamma(0.1,0.1)
  }
  tau[5] ~ dgamma(0.1,0.1)

  for(age in 1:30){for(j in 1:4){
    PART[age,j] <- (log(age)-c[j])/d[j]
    fitted[age,j] <- a[j] + b[j]/(1+exp(-PART[age,j]))
  }}
}")

model4 <- jags.model(model_string4,data = data,quiet=TRUE, n.chains=2)
update(model4, burn, progress.bar="none")

samples4 <- coda.samples(model4, variable.names=c("a","b","c","d","fitted","tau"), n.iter=iters, thin=thin, pr

ESS <- effectiveSize(samples4)
ESS[which.min(ESS)]

```

```

##      c[3]
## 303.9109

```

```
ESS[which.max(ESS)]
```

```

## fitted[2,1]
##      19001.98

```

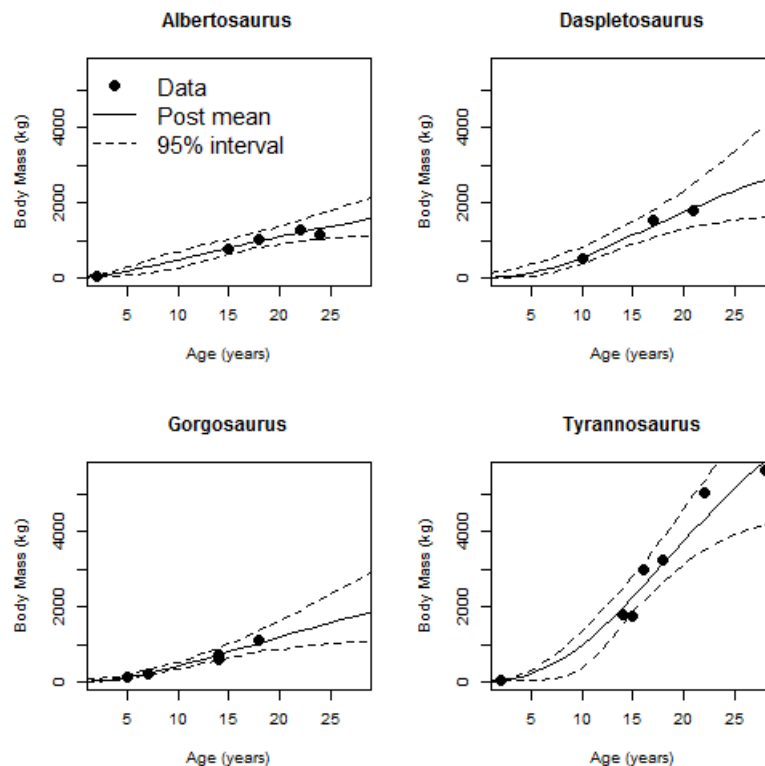
```

fit4 <- summary(samples4)$quantiles[1:120+16,3]
lo4 <- summary(samples4)$quantiles[1:120+16,1]
hi4 <- summary(samples4)$quantiles[1:120+16,5]
id <- rep(1:4,each=30)

par(mfrow=c(2,2))
for(j in 1:4){
  plot(NA,xlim=range(age),ylim=range(mass),
       xlab="Age (years)", ylab="Body Mass (kg)",
       main=names[j])
  points(age[sp==j],mass[sp==j],pch=19,cex=1.5)
  lines(1:30,fit4[id==j],lty=1)
  lines(1:30,lo4[id==j],lty=2)
  lines(1:30,hi4[id==j],lty=2)

  if(j==1){
    legend("topleft",c("Data", "Post mean", "95% interval"),
          pch=c(19,NA,NA),lty=c(NA,1,2),cex=1.5,bty="n")
  }
}

```



Summary: Compared to Model 3, this fit has better convergence and much more narrow 95% intervals. By sharing information across curves we are able to estimate non-linear means for each species.

Model comparison

```

# No pooling, log-linear model
dic1 <- dic.samples(model1, n.iter=iters, thin=thin, progress.bar="none")
dic1

```

```

## Mean deviance: 3.496
## penalty 25
## Penalized deviance: 28.5

```

```

# Pooling, log-linear model
dic2 <- dic.samples(model2, n.iter=iters, thin=thin, progress.bar="none")
dic2

```

```
## Mean deviance: -12.54
## penalty 9.154
## Penalized deviance: -3.386
```

```
# No pooling, logistic model
dic3 <- dic.samples(model3, n.iter=iters, thin=thin, progress.bar="none")
dic3
```

```
## Mean deviance: 22.67
## penalty 41.59
## Penalized deviance: 64.26
```

```
# Pooling, logistic model
dic4 <- dic.samples(model4, n.iter=iters, thin=thin, progress.bar="none")
dic4
```

```
## Mean deviance: -14.02
## penalty 12.2
## Penalized deviance: -1.828
```

Summary: Both hierarchical models that borrow information across taxa have much smaller DIC than models that analyze taxa separately. The log-linear model edges out the logistic model for the smallest DIC (best model).

Processing math: 100%